



Neural Networks in the Wild

HANDWRITING RECOGNITION

BY JOHN LIU

Motivation

- US Post Office (700 million pieces of mail per day)
 - HWAI (Lockheed-Martin 1997)
 - Letter Recognition Improvement Program
- Electronic Health Records (*Medical scribble*)
- Banking: check processing
- Legal: signature verification
- Education: autograding



OCR is not HWR

Optical Character Recognition

Handwriting Recognition

Fixed fonts

Free flowing

No character overlap

Overlapping characters

Easy alignment id

Flowing alignment

Fixed aspect ratios


Variable aspect ratios

Low noise

Can be noisy

HWR: Online vs Offline


- Online recognition = conversion of text as it is written
 - Palm PDA, Google Handwrite
 - Low noise = easier classification
 - Features: stroke pressure, velocity, trajectory
- Offline recognition = image conversion of text
 - Computer-vision based methods to extract glyphs
 - Much more difficult than normal OCR



Offline HWR

- Preprocessing
 - Discard irrelevant artifacts and remove noise
 - Smoothing, thresholding, morphological filtering
- Segmentation and Extraction
 - Find contours and bounding boxes
 - Glyph extraction
- Classification
 - Linear and ensemble methods
 - Neural Networks

HWR Examples



Scanned handwritten note

- noisy background
- varying character size
- biased ground truth

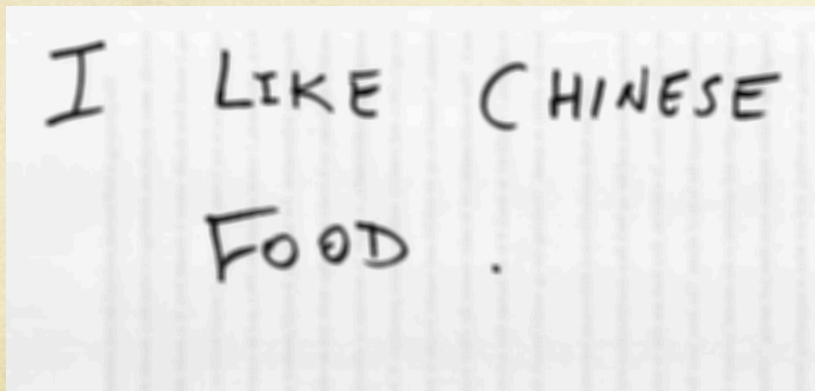
License Plate image

- goal: plate ID
- similar problem to HWR
- registration renewal soon



Kernel Smoothing

- Blurring helps eliminate noise after thresholding
 - Local kernel defines averaging area
 - Used 8x8 kernel for example images
 - OpenCV: `cv2.blur()`



Thresholding

- Thresholding converts to b/w image
 - Colors are inverted so (black,white) = (0,1)
 - Some noise is present, but greatly reduced due to smoothing in previous step
 - OpenCV: `cv2.threshold()`



Morphological Filtering

- Erosion: eat away at the boundaries of objects
 - Removes white noise and small artifacts
 - OpenCV: `cv2.erode()` with 4x4 kernel
- Dilation: increases thickness and white region
 - opposite of erosion
 - Useful in joining broken parts of an object
 - OpenCV: `cv2.dilate()` with 2x2 kernel

Opening vs Closing

- Erosion/Dilation = Opening
 - Eliminates noise outside objects

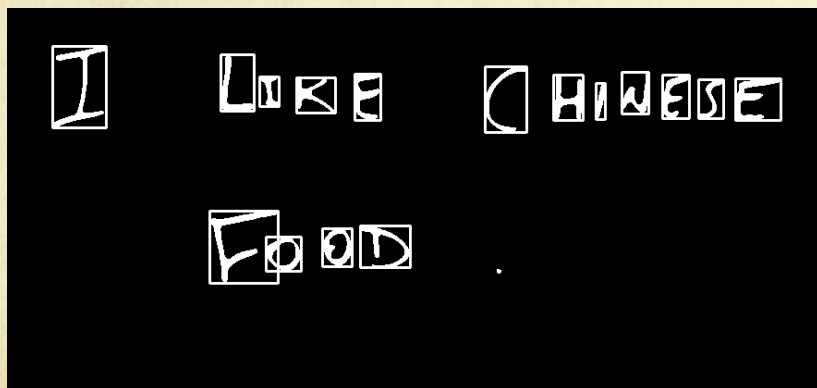


- Dilation/Erosion = Closing
 - Eliminates noise within objects



Contour/Bounding Box

- Contours = set of all outer contiguous points
 - Approximate contours as a reduced polygon
 - Calculate the bounding rectangle
 - OpenCV: `cv2.findContours()`, `cv2.approxPolyDP()`, `cv2.boundingRect()`



Glyph Extraction

- Results from bounding boxes



I L I K E

C H W N E S E

F O O D

NIST Special DB 19

- Contains 814,255 segmented handwritten characters
- Superset of MNIST that includes alphabetic characters
- 62 character classes [A-Z], [a-z], [0-9], 128x128 pixels
- We down-sample to 32x32 and use only a subsample of 90,000 characters (train=70,000, test&valid=10,000)






MNIST vs SD19

MNIST (LeCun)	SD19
10 classes	62 classes
Digits	Upper & Lower case + Digits
28x28 pixel	128x128 pixel
60,000 samples	814,255 samples
boring	GROOVY

Classification

- Goal: correctly classify character with highest:
 - Accuracy
 - $F1$ = geometric mean of precision & recall
- Typical Methods
 - Linear
 - SVM
 - Ensemble
 - Neural Networks



Logistic Regression

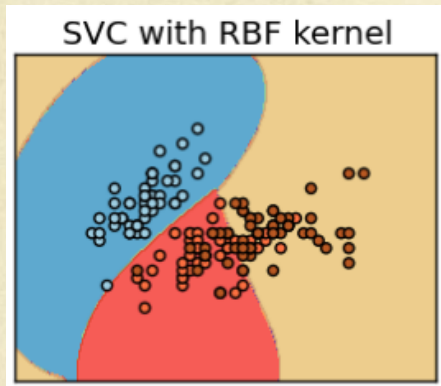
- a.k.a. Softmax, MaxEnt, logit regression
- We use multinomial LR with classes = 62
- Implemented with scikit-learn

```
modelSG = linear_model.SGDClassifier(loss='log',penalty='l2',alpha=0.001,  
                                     n_iter=1,shuffle=False,n_jobs=-1,random_state=33)  
modelSG.fit(trainX,trainY)
```

Accuracy score = 59%, avg F1 score = 0.56 (baseline)

SVM - RBF

- We use a Gaussian kernel (a.k.a. Radial Basis Function)
- Implemented with scikit-learn (slow!)



```
modelSVC = svm.SVC(kernel='rbf', gamma=0, tol=0.01,  
                    verbose=False, max_iter=1, random_state=33)  
modelSVC.fit(trainX, trainY)
```

Accuracy score = 65%, avg F1 score = 0.61

Random Forest

- Ensemble estimator that builds a random forest of decision trees and combines their estimations
- Crowd Intelligence: “the wisdom of Crowds”
- Implemented with scikit-learn (fast!)

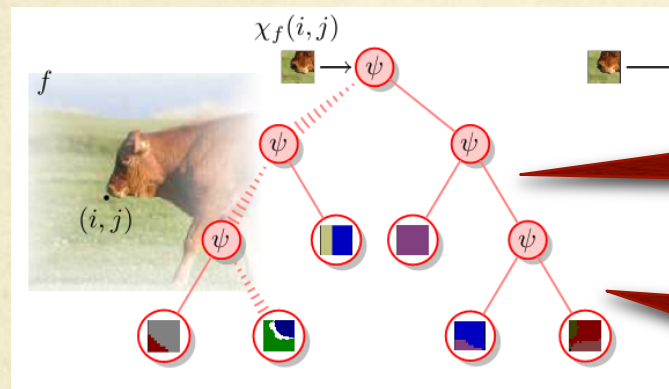


```
modelRF = RandomForestClassifier(n_estimators=10,criterion='gini',  
                                max_depth=None,n_jobs=-1,verbose=0,random_state=33)  
modelRF.fit(trainX,trainY)
```

Accuracy score = 69%, avg F1 score = 0.66

Extra Trees

- a.k.a. Extremely Randomized Trees, similar to Random Forest except splits are also randomized



RF picks best from random subset of features

Extra Trees picks randomly from random subset of features


```
modelET = ExtraTreesClassifier(n_estimators=10,  
                               max_depth=None, random_state=33)  
modelET.fit(trainX, trainY)
```

Accuracy score = 73%, avg F1 score = 0.71

Neural Networks

- HWR inherently a computer vision problem, apt for neural networks given recent advances
- Inputs: image reshaped as a binary vector
- Outputs: one-hot representation = 62 bit vector
- Question: How do you actually go about building a (deep) neural network model?

WARNING: don't try this in Excel



Theano

- Python library that implements Tensor objects that leverage GPUs for calculation
 - Plays well with `numpy.ndarray` datatypes
 - Transparent use of GPU (float32) = 140X faster
 - Dynamic C code generation
- Installation is not fun:
 - Update GPU driver
 - Install CUDA 6.5 + Toolkit
 - Install CUDAMat

PyLearn2

- Python machine learning library toolbox built upon Theano (GPU speed)
- Provides flexibility to build customized neural network models with full access to hyper-params (nodes per layer, learning rate, etc...)
- Uses pickle for file I/O
- Two methods to implement models:
 - YAML
 - ipython notebook



Pylearn2 Code Structure


- Dataset specification
 - Uses DenseDesignMatrix class
- Model configuration
 - LR, Kmeans, Softmax, RectLinear, Autoencoder, RBM, DBM, Cuda-convnet, more...
- Training Algorithm choice
 - SGD, BGD, Dropout, Corruption
- Training execution
 - Train class
- Prediction and results

#1 Feedforward NN

- Hidden layer = sigmoid, Output layer = softmax



Input = 1024
(=32x32) bit
vector



Prediction = 62
bit vector
(one-hot rep)

Hyperparams

- Hidden layer = 400 neurons
- Output layer = 62 neurons
- Random initialization (symmetry breaking)
- SGD algorithm (mini-batch stochastic gradient descent)
 - Fixed learning rate = 0.05
 - Batch size = 100
- Termination = after 100 epochs

#1 pylearn2 code

```
h0 = mlp.Sigmoid(layer_name="h0",dim=400, sparse_init=20)
y0 = mlp.Softmax(n_classes=62,layer_name="y0",sparse_init=20)
layers = [h0, y0]

nn = mlp.MLP(layers,nvis=1024)
algo = sgd.SGD(learning_rate=0.05,batch_size=100,monitoring_dataset=valid,
               termination_criterion=EpochCounter(100))
algo.setup(nn,train)

save_best = best_params.MonitorBasedSaveBest(channel_name="y0_misclass",
                                              save_path='best_params.pkl')

while True:
    algo.train(dataset=train)
    nn.monitor.report_epoch()
    nn.monitor()
    save_best.on_monitor(nn,train,algo)
    if not algo.continue_learning(nn):
        break
```


#1 Running in ipython

Epochs seen: 100

Batches seen: 70000

Examples seen: 7000000

learning_rate: 0.0500000119209

objective: 0.623726010323

y0_col_norms_max: 11.5188903809

y0_col_norms_mean: 6.90935611725

y0_col_norms_min: 4.80359125137

y0_max_max_class: 0.996758043766

y0_mean_max_class: 0.759563267231

y0_min_max_class: 0.221193775535

y0_misclass: 0.173800006509

y0_nll: 0.623726010323

y0_row_norms_max: 6.41655635834

y0_row_norms_mean: 2.55642676353

y0_row_norms_min: 0.48346811533

Accuracy = 83%


Problem: Overfitting

- It's easy to overfit using neural networks
 - How many neurons per layer? 400? 800?
- Methods to deal with it include:
 - L1, L2, ElasticNet regularization
 - Early stopping
 - Model averaging
 - **DROPOUT**



Remedy: Dropout

- Dropout invented by G. Hinton to address overfitting
 - Automatically provides ensemble boosting




- Prevents neurons from co-adapting

#2 NN w/ Dropout

- 2 Softplus hidden layers, Softmax output layer




Input = 1024
(=32x32) bit
vector



S

Prediction = 62
bit vector
(one-hot rep)

$$\text{Softplus } f(x) = \log(1 + e^x)$$



#2 pylearn2 code

```
# SoftPlus with Dropout

h0 = mlp.Softplus(layer_name='h0', dim=800, sparse_init=40)
h1 = mlp.Softplus(layer_name='h1', dim=800, sparse_init=40)
y0 = mlp.Softmax(layer_name='y0', n_classes=62, irange=0)
layers = [h0, h1, y0]

model = mlp.MLP(layers, nvis=1024)

monitoring = dict(valid=valid)
termination = MonitorBased(channel_name="valid_y0_misclass", N=10)
extensions = [best_params.MonitorBasedSaveBest(channel_name="valid_y0_misclass",
save_path="train_best.pkl")]

algorithm = sgd.SGD(0.1, batch_size=100, cost=Dropout(),
                    monitoring_dataset = monitoring,
                    termination_criterion = termination)

print 'Running training'
train_job = Train(train, model, algorithm, extensions=extensions,
                  save_path="train.pkl", save_freq=1)
train_job.main_loop()
```

Termination Condition = stop if no improvement after N=10 epochs

#2 running in ipython

Epochs seen: 93

Batches seen: 65100

Examples seen: 6510000

learning_rate: 0.1000000023842

total_seconds_last_epoch: 10.1289653778

training_seconds_this_epoch: 6.20919704437

valid_objective: 1.4562972784

Accuracy = 76%

valid_y0_col_norms_max: 5.99511814117

valid_y0_col_norms_mean: 2.90585327148

valid_y0_col_norms_min: 2.15357899666

valid_y0_max_max_class: 0.967477440834

valid_y0_mean_max_class: 0.583315730095

valid_y0_min_max_class: 0.12644392252

valid_y0_misclass: 0.253300011158

valid_y0_nll: 0.946564733982


valid_y0_row_norms_max: 2.06650829315

valid_y0_row_norms_mean: 0.773801326752

valid_y0_row_norms_min: 0.347277522087


Problem: SGD speed

- SGD learning tends to be slow when curvature differs among features
- Errors change trajectory of gradient, but always perpendicular to feature surface



Remedy: Momentum

- Errors change velocity of gradient, not gradient itself




- Start with small momentum to dampen oscillations
- Fully implemented in Pylearn2
- Other methods (conjugate gradient, Hessian-free)

#3 NN w/ Momentum

- Rectified Linear hidden layer, Softmax output layer



Input = 1024
(=32x32) bit
vector




S

Prediction= 62
bit vector
(one-hot rep)

Hidden = 400 neurons
Initial Momentum = 0.5
Final Momentum = 0.99

Rectified Linear
 $f(x) = \max(x, 0)$



#3 pylearn2 code

```
# Rectified Linear with Momentum

from pylearn2.training_algorithms import sgd, learning_rule

h0 = mlp.RectifiedLinear(layer_name='h0', dim=400, sparse_init=40)
y0 = mlp.Softmax(layer_name='y0', n_classes=62, irange=0)
layers = [h0, y0]

model = mlp.MLP(layers, nvis=1024)

# momentum
initial_momentum = 0.5
final_momentum = 0.99
start = 1
saturate = 50
momentum_rule = learning_rule.Momentum(initial_momentum)

monitoring = dict(valid=valid)
termination = MonitorBased(channel_name="valid_y0_misclass", N=10)
extensions = [best_params.MonitorBasedSaveBest(channel_name="valid_y0_misclass",
                                                save_path="rect_best.pkl"),
              learning_rule.MomentumAdjustor(final_momentum, start, saturate)]

algorithm = sgd.SGD(0.1, batch_size=100, cost=Dropout(), learning_rule=momentum_rule,
                    monitoring_dataset = monitoring, termination_criterion = termination)

print 'Running training'
train_job = Train(train, model, algorithm, extensions=extensions,
                  save_path="rect.pkl", save_freq=5)
train_job.main_loop()
```



#3 running in ipython

Epochs seen: 27
Batches seen: 18900
Examples seen: 1890000
learning_rate: 0.100000023842
momentum: 0.760000526905
total_seconds_last_epoch: 4.66804122925
training_seconds_this_epoch: 2.96171355247
valid_objective: 1.42965459824
valid_y0_col_norms_max: 4.5757818222
valid_y0_col_norms_mean: 3.47835850716
valid_y0_col_norms_min: 2.6680624485
valid_y0_max_max_class: 0.962530076504
valid_y0_mean_max_class: 0.555216372013
valid_y0_min_max_class: 0.12590457499
valid_y0_misclass: 0.249799996614
valid_y0_nll: 0.971761405468
valid_y0_row_norms_max: 2.2137401104
valid_y0_row_norms_mean: 1.3540699482
valid_y0_row_norms_min: 0.666570603848

Accuracy = 76%

Autoencoders

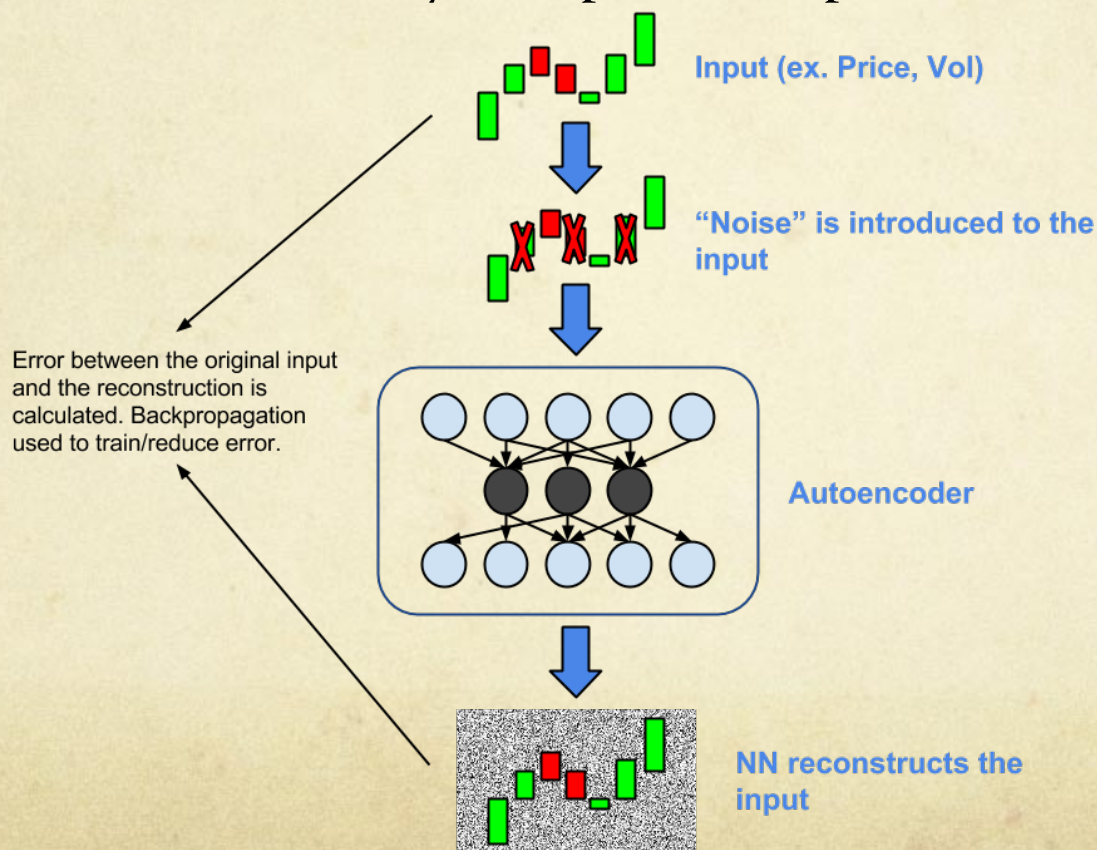
- Learns efficient codings (dimensionality reduction)
 - Linear units = similar to PCA (+rotation)
 - Nonlinear units = manifold learning



- Autoencoders are deterministic (not good at predicting)

Denoising Autoencoders


- Corrupt input with noise (randomly set to 0), then force hidden layer to predict input



#4 Stacked DAE

- Autoencoders for hidden layers 1 & 2, softmax output layer

Unsupervised
Training for
Autoencoder
Hidden Layers



Supervised
Training to
fine-tune
Stacked DAE

- Pre-train Hidden Layers 1 & 2 sequentially, stack a softmax output layer on top and fine tune

#4 training layer 1

- Pre-train Layer 1, save weights

```
#layer 1

corruptor = BinomialCorruptor(corruption_level=0.2)
model = autoencoder.DenoisingAutoencoder(nvis=1024,nhid=500,irange=0.05,
                                         corruptor=corruptor,act_enc="tanh",act_dec=None)
algorithm = sgd.SGD(learning_rate=0.001, batch_size=100,
                   monitoring_batches=5, cost=MeanSquaredReconstructionError(),
                   monitoring_dataset=train, termination_criterion=EpochCounter(10))

print 'Running training'
train_job = Train(train, model, algorithm, save_path="dae_layer1.pkl", save_freq=1)
train_job.main_loop()
```

Epochs seen: 10

Batches seen: 7000

Examples seen: 700000

learning_rate: 0.00100000000475

objective: 9.47059440613

total_seconds_last_epoch: 5.54894971848

training_seconds_this_epoch: 3.88439941406

#4 training layer 2

- Pre-train layer 2, using outputs of layer 1 as inputs

```
# layer 2

transformer = serial.load("dae_layer1.pkl")
ntrain = TransformerDataset(raw=train,transformer=transformer)

corruptor = BinomialCorruptor(corruption_level=0.3)
model = autoencoder.DenoisingAutoencoder(nvis=500,nhid=500,irange=0.05,
                                         corruptor=corruptor,act_enc="tanh",act_dec=None)
algorithm = sgd.SGD(learning_rate=0.001, batch_size=100,
                    monitoring_batches=5, cost=MeanSquaredReconstructionError(),
                    monitoring_dataset=ntrain, termination_criterion=EpochCounter(10))

print 'Running training'
train_job = Train(ntrain, model, algorithm, save_path="dae_layer2.pkl", save_freq=1)
train_job.main_loop()
```

Epochs seen: 10

Batches seen: 7000

Examples seen: 700000

learning_rate: 0.00100000000475

objective: 2.91170930862

total_seconds_last_epoch: 4.94157600403

training_seconds_this_epoch: 3.63193702698

#4 Fine-tuning

○ SGD with momentum for 50 Epochs

```
# stacking and supervised fine-tuning

dae1 = serial.load("dae_layer1.pkl")
dae2 = serial.load("dae_layer2.pkl")

h1 = mlp.PretrainedLayer(layer_name='h1', layer_content=dae1)
h2 = mlp.PretrainedLayer(layer_name='h2', layer_content=dae2)
y0 = mlp.Softmax(layer_name='y0', n_classes=62, irange=0.005, max_col_norm=1.9365)
layers = [h1, h2, y0]

monitoring = dict(valid=valid)
callback = sgd.ExponentialDecay(decay_factor=1.00004, min_lr=0.000001)
extensions = [best_params.MonitorBasedSaveBest(channel_name="valid_y0_misclass",
                                                save_path="dae_best.pkl"),
              learning_rule.MomentumAdjustor(final_momentum=0.7, start=1, saturate=250)]
model = mlp.MLP(layers, batch_size=100, nvis=1024)
algorithm = sgd.SGD(learning_rate=0.05, init_momentum=0.5, batch_size=100,
                    update_callbacks=callback,
                    monitoring_dataset=monitoring, termination_criterion=EpochCounter(50))

print 'Running training'
train_job = Train(train, model, algorithm, extensions, save_path="dae.pkl", save_freq=1)
train_job.main_loop()
```

Classification Results

Model	Accuracy
Logistic Regression	59%
SVM w/RBF	65%
Random Forest	69%
Extremely Random Forest	73%
2-layer Sigmoid	83%
3-layer Softplus w/dropout	76%
2-layer RectLin w/momentum	76%
Stacked DAE w/momentum	89%

Tools

- OpenCV – computer vision library
- Scikit-image – image processing library
- Scikit-learn – ML on python
- Theano – fast tensor math library
- Pylearn2 – neural networks on python
- Nolearn – deep learning on python
- DL4J – deep learning on Java
- Torch7 – ML on LUA
- MLlib – ML on Spark

LAST WORDS

I LIKE CHINESE FOOD

- STACKED DAE